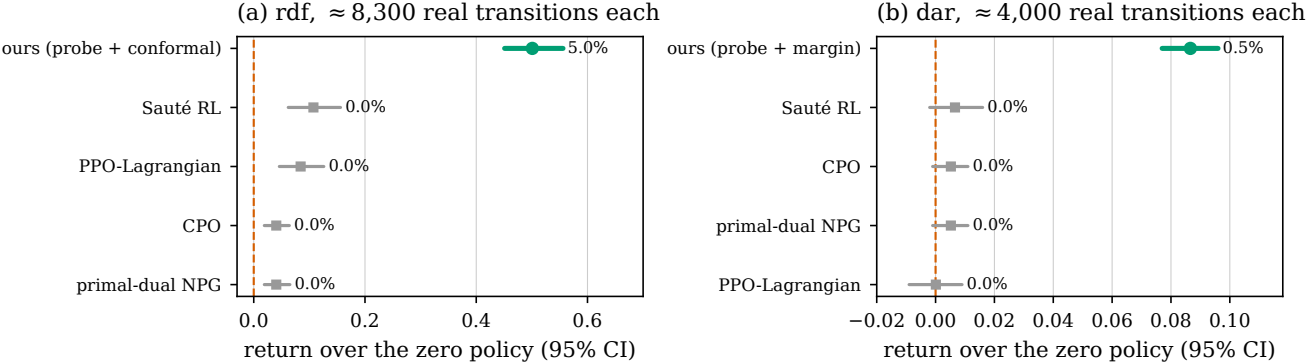
(a) rdf, ≈ 8,300 real transitions each
ours (probe + conformal)
Sauté RL
PPO-Lagrangian
CPO
primal-dual NPG
5.0%
0.0%
0.0%
0.0%
0.0%
0.0
0.2
0.4
0.6
return over the zero policy (95% CI)
(b) dar, ≈ 4,000 real transitions each
ours (probe + margin)
Sauté RL
CPO
primal-dual NPG
PPO-Lagrangian
0.5%
0.0%
0.0%
0.0%
0.0%
−0.02
0.00
0.02
0.04
0.06
0.08
0.10
return over the zero policy (95% CI)