

EXPLAIN

$$F(b) = \exp(-\text{KL}(\pi_\phi \parallel \hat{\pi}_b)/K)$$

WHAT IT DOES

A brief is parsed by a frozen deterministic parser back into an action distribution, then scored against the policy.

Split conformal: with calibration scores $s_i = 1 - F(b_i)$, the $\lceil (n+1)(1-\delta) \rceil$ order statistic bounds the floor.

Per dimension KL is the only comparable form; summed KL sends every score to zero above about 16 actions.

WHAT HELD UP

Certificate coverage holds at 90 and 95 percent on 3 of 3 seeds on real fire plans, with a floor of 0.706 that is not vacuous. It reports honestly on a weak generator rather than overstating it.

WHAT DID NOT

The faithfulness claim is withdrawn. A permutation control scores each brief against a different state's policy and finds a gap of zero, the signature of a generator that ignores its condition.

Three fixes were tried and all were negative: wider conditioning, a differentiable contrastive objective, and 2.5 times the training budget. The testbed also could not have answered the question, since its trained policy varies 1.2 percent and yields 8 distinct actions across 64 states.